

Variational Inference - Assignment 2

Printout

Ravi Regulatedda

README

To run the code, just run the code.py file. You will be asked for a choice of dataset 1. 500-500-30 and 2. 100-k. After making this choice, part 1 and part 2 is printed to stdout and the plots are created as png files in the directory.

Code

```
# homework.py
from scipy.sparse import lil_matrix
from matrix import *
import scipy.sparse
import time
import matplotlib.pyplot as plt

print("Please enter choice of dataset: \n")
print("1. 500-500-3\n")
print("2. 100-k\n")
choice = int(input("Please enter your choice: 1 (or) 2 [To Debug, type 0] - "))

name = "40-40-2"
data = "hw2data/data-40-20-2.txt"
rank = [1]
rank1 = 1

if choice == 1:
    name = "500-500-3"
    data = "hw2data/data-500-500-3.txt"
    rank = [1, 2, 3, 5, 10, 20]
    rank1 = 3

if choice == 2:
    data = "hw2data/ml-100k.txt"
    rank = [1, 2, 3, 5, 20, 100]
    rank1 = 5

# reading the dataset and storing in a sparse matrix

with open(data, encoding="utf-8") as d:
    test = []
    instance = d.readline()
    # splitting the instance at space
    users, items, scores = map(int, instance.split())
    # storing users and items into a list of lists (sparse) matrix
    m = scipy.sparse.lil_matrix((users, items))
    # doing a half split
    split = int(scores / 2)
    for _ in range(split):
        inst = d.readline().split()
        x = int(inst[0])
        y = int(inst[1])
        # m will be our training data
        m[x, y] = float(inst[2])
    for _ in range(scores - split):
        inst = d.readline().split()
        x = int(inst[0])
```

```

        y = int(inst[1])
        test.append((x, y, float(inst[2])))

# making sure m is a sparse matrix
m = m.tocsr()

# baseline
a, b = [], []
c = scipy.sparse.find(m)[2].mean()
for i in range(0, m.shape[0]):
    if len(scipy.sparse.find(m[i])[0]) == 0:
        a.append(c)
    else:
        a.append(m[i].mean())
for j in range(0, m.shape[1]):
    if len(scipy.sparse.find(m.T[j])[0]) == 0:
        b.append(c)
    else:
        b.append(m.T[j].mean())

# defining an RMSE function (since we have to use it many times)
def rmse(true, pred):
    # convert them to np.array just in case
    true = np.array(true)
    pred = np.array(pred)
    difference_2 = (true - pred) ** 2
    error = np.sqrt(difference_2.mean())
    return error

# calculating the baseline score and the baseline rmse
true_scores, base_preds = [], []
for instance in test:
    # unpacking
    user, item, score = instance
    true_scores.append(score)
    base_preds.append((a[user] + b[item]) / 2)
    base_rmse = rmse(true_scores, base_preds)

print("Baseline rmse is: ", base_rmse)

# part 1 (given dataset/rank has already been chosen)
part1_rmse, times = [], []
start = time.time()
time_taken = 0
# since we used a generator, we can step through each u, v and see how the rmse
# goes down per iteration (till convergence or end of 100 iterations)
for u, v in matrix_factorization(m, rank1):
    part1_preds = []
    end = time.time()
    time_taken += end - start
    times.append(time_taken)
    for instance in test:
        user, item, score = instance
        pred = u[user].dot(v[item]) # making the predictions
        part1_preds.append(pred)
    part1_rmse.append(rmse(true_scores, part1_preds)) # rmse per iteration
    start = time.time() # resetting time for next iteration

print("For Part 1: \n")
print("Total Iterations: ", len(part1_preds), "\n")
if len(part1_preds) < 100:
    print("We had early stopping due to difference between iterations < 0.01\n")
print("RMSE: ", part1_rmse)
print("Time taken: ", times)
print("\n")

# part 2
# part 2 (given dataset has already been chosen)
part2_rmse = []
for r in rank:
    # we do this to get the last item from our generator

```

```

output = list(matrix_factorization(m, r))
final = output[-1]
u, v = final
part2_preds = []
for instance in test:
    user, item, score = instance
    pred = u[user].dot(v[item])
    part2_preds.append(pred)
part2_rmse.append(rmse(true_scores, part2_preds))

print("For Part 2: ")
print("For the given ranks: ", rank)
print("The RMSE changes as follows: ", part2_rmse)

plt.plot(times, part1_rmse)
plt.title(name + " Part 1")
plt.xlabel("Time taken ")
plt.ylabel("RMSE")
plt.savefig("Part 1 " + name + ".png")

plt.clf()

plt.plot(rank, part2_rmse)
plt.xticks(rank)
plt.title(name + " Part 2")
plt.xlabel("Rank")
plt.ylabel("Change in RMSE")
plt.savefig("Part 2 " + name + ".png")

print("Please check your current directory for generated plots")

# matrix.py
import numpy as np
from scipy.sparse import csr_matrix
import scipy.sparse

# defining the matrix decomposition
# this is a generator i.e, it 'yields' a value after 100 iterations
# that makes it easy for us to store and plot the performance metrics
# i used black to format my code. if it looks strange blame that :)
def matrix_factorization(m, rank):
    """Performs matrix factorization according to Lim & Teh, 2007. Takes in m: a sparse matrix
    and rank"""
    # initialize the variables
    # shape of the matrix
    # this assumes we have a 2D matrix (which we do)

    i_m = m.shape[0]
    j_m = m.shape[1]
    d = rank

    # variables and constants
    sigma_2 = np.ones(d)
    rho_2 = np.ones(d) / d
    tau_2 = 1
    u = []
    v = []
    # notation from paper
    S, phi, psi, t = [], [], [], []

    # initializing our variables with ones and random values according to the question
    for i in range(i_m):
        phi.append(np.eye(d))
        u.append(np.random.normal(0, 1, d))
    for j in range(j_m):
        psi.append(np.eye(d))
        S.append(np.diag(1 / rho_2))
        t.append(np.zeros(d))
        v.append(np.random.normal(0, 1, d))

    # make sure everything is an numpy array so there are no runtime/indexing errors
    phi = np.array(phi)

```

```

psi = np.array(psi)
S = np.array(S)
t = np.array(t)
u = np.array(u)
v = np.array(v)

norm_u = 0
norm_v = 0

N = []
for i in range(i_m):
    # from the paper N[i] is 1 if we observe m_ij. Since we observe a 0 or 1, we just
    # append it
    N.append(scipy.sparse.find(m[i])[1])
# all the places in the sparse matrix where there are elements
ob = scipy.sparse.find(m)

# we perform our EM now
# iterating 100 times
for _ in range(100):
    # E step
    # Q(U)
    for i in range(0, i_m):
        # this stores the matrix product (outer product in numpy)
        outer = np.zeros((d, d))
        N_i = N[i]
        for j in N_i:
            outer += np.outer(v[j], v[j])
        phi[i] = np.linalg.inv(
            np.diag(1 / sigma_2) + (psi[N_i].sum(0) + outer) / tau_2
        )
        mtplr = ((m[i, N_i] * (v[N_i])) / tau_2).sum(0)
        u[i] = phi[i].dot(mtplr)
        S[N_i] += (phi[i] + np.outer(u[i], u[i])) / tau_2
        t[N_i] += np.outer(csr_matrix.toarray(m[i, N_i]), (u[i])) / tau_2

    # Q(v)
    psi = np.linalg.inv(S)
    for j in range(j_m):
        v[j] = psi[j].dot(t[j])

    # M step
    for l in range(d):
        sigma_2[l] = ((phi[:, l, l] + u[:, l] ** 2).sum()) / (i_m - 1)

K = len(ob[1])
Tr = 0 # trace
for i, j in np.array([ob[0], ob[1]]).T:
    A = phi[i] + np.outer(u[i], u[i])
    B = psi[j] + np.outer(v[j], v[j])
    Tr += np.trace(A.dot(B))

tau_2 = (
    (
        (ob[2] ** 2) - (2 * ob[2] * np.einsum("ij,ij->i", u[ob[0]], v[ob[1]]))
    ).sum()
    + Tr
) / (K - 1)

new_norm_u = np.linalg.norm(u)
new_norm_v = np.linalg.norm(v)

if abs(new_norm_u - norm_u) < 0.01 or abs(new_norm_v - norm_v) < 0.01:
    # early stopping
    break
else:
    norm_u, norm_v = new_norm_u, new_norm_v
yield np.array(u), np.array(v)

```

Outputs

```
> python homework.py
Please enter choice of dataset:
1. 500-500-3
2. 100-k

Please enter your choice: 1 (or) 2 [To Debug, type 0] - 1
Baseline rmse is: 1.8019448325127836
For Part 1:

Total Iterations: 25000

RMSE: [1.781923883557573, 1.5937212241838992, 1.2321690119758026, 0.722144708569602, 0.3283120448423085, 0.23598144269743926, 0.21916363355676868, 0.2160546121083858]
Time taken: [0.48094797134399414, 0.8806238174438477, 1.3032078742980957, 1.7859125900268555, 2.105543613433838, 2.519871717380957, 2.923938751220783, 3.3263778686523438]

For Part 2:
For the given ranks: [1, 2, 3, 5, 10, 20]
The RMSE changes as follows: [1.5068963472883776, 1.1600153878586528, 0.21661693422373104, 0.22146709229941794, 0.24548545758228488, 0.2738502473472803]
Please check your current directory for generated plots
```

Figure 1: Output after running matrix factorization on 500-500 dataset

```
> python homework.py
Please enter choice of dataset:
1. 500-500-3
2. 100-k

Please enter your choice: 1 (or) 2 [To Debug, type 0] - 2
Baseline rmse is: 3.4397503673365066
For Part 1:

Total Iterations: 50000

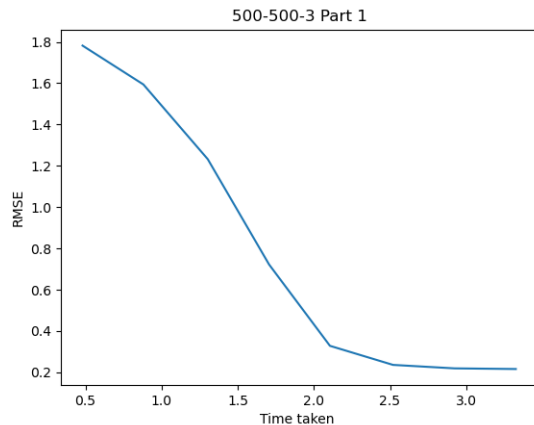
RMSE: [3.4112381981697038, 1.7459647590787626, 1.259336614373994, 1.0717570365475515, 1.0094145671801862, 0.9869535940433072, 0.9773663235991397, 0.9729829997699876, 0.971110051528889, 0.9706030114922143, 0.970895231935521, 0.9716668351359257, 0.9727191894855078, 0.973922483582565, 0.9751909066372002, 0.9764687906820774, 0.9777215207516838, 0.9789289840670536, 0.9800807532549812, 0.9811726218729699, 0.9822042020987551, 0.9831773211356803, 0.9840949888228743, 0.9849607561742121, 0.9857783316682548, 0.9865513619941978, 0.9872833146365183, 0.9879774217741703, 0.9886366600621885, 0.9892637507258457, 0.9898611706150541, 0.9904311686869036, 0.99095757846923342, 0.9914968682259249, 0.9919960971235808, 0.9924749946841002, 0.9929349454815756, 0.9933772097052951, 0.9938029360595462, 0.9942131733075552, 0.9946088805674178, 0.9949308364767786, 0.9953601472413681, 0.9957172543769532, 0.9960629401452848, 0.9963978342757025, 0.9967225185542518, 0.9970375314535498, 0.9973433721682833, 0.9976405042139241, 0.9979293586393044, 0.9982103368982594, 0.9984838134197467, 0.9987501379115318, 0.9990096374283061]
Time taken: [1.014092206954956, 1.9421393871307373, 2.8472652435302734, 3.7813050746917725, 4.687999963760376, 5.588575124740601, 6.562986135482788, 7.526334047317505, 8.519048929214478, 9.499677896499634, 10.39844298362732, 11.366978168487549, 12.267967224121094, 13.193376302719116, 14.053032159805298, 14.950924396514893, 16.00252628326416, 16.864455461502075, 17.84455132484436, 18.75015616416931, 19.7137610912323, 20.627712965011597, 21.552350997924805, 22.5420880317688, 23.435638189315796, 24.443799257278442, 25.36793327331543, 26.301979303359985, 27.26687240600586, 28.163235425949097, 29.044275283813477, 29.99522614479065, 30.884100437164307, 31.81834840774536, 32.71967530250549, 33.687013149261475, 34.61783504486084, 35.56900691986084, 36.43928599357605, 37.2837769985199, 38.5317702293396, 39.508172273635864, 40.4709312915802, 41.342119216918945, 42.29035234451294, 43.22833251953125, 44.223445415496826, 45.111032247543335, 46.00450420379639, 46.965280294418335, 47.989593267440796, 48.999062299728394, 50.0300931930542, 51.08395743370056, 52.17591118812561]

For Part 2:
For the given ranks: [1, 2, 3, 5, 20, 100]
The RMSE changes as follows: [0.9709818673837846, 0.9801935787643553, 0.9847149864732342, 1.0167913768950378, 1.1632308654873804, 1.1061426231943978]
Please check your current directory for generated plots
```


Figure 2: Output after running matrix factorization on 100-k dataset

Results

From the graphs below, we can see that the RMSE for both the datasets goes down per iteration, showing that they converge to a minimum.

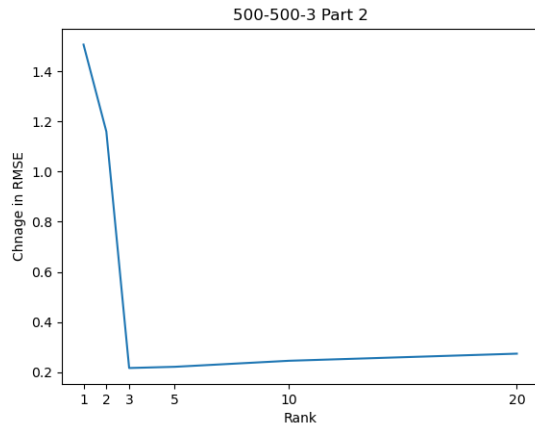


(a) Time v/s RMSE for 500-500 dataset




(b) Time v/s RMSE for 100k dataset

For the $500 - 500 - 3$ dataset, the RMSE falls sharply at rank 2 and 3 from rank 1. Following that the RMSE plateaus and rises only slightly. Meanwhile, in the 100k dataset, lowest RMSE is at rank 1 and rises sharply till rank 20 following which it falls as we move towards higher rank but not as low.



(a) Rank v/s RMSE for 500-500 dataset



(b) Rank v/s RMSE for 100k dataset

For both the datasets, the performance is hugely better than the baseline.